



# Which Mushroom Is That?

Thanks to Springboard mentor



Branko Kovac, Logikka

David Olivero

April 27<sup>th</sup> 2020 cohort



Modeling accomplished with:




# The Issue

- Mushroom hunting is rewarding, but daunting
- Huge variety of mushrooms
- Some are delicious
- Some are deadly
- How can you tell what you're looking at?



# The Question

If you wanted some help identifying mushroom genus (for example, from an app on your smart phone), would it be possible?



Can a model be built to properly classify a mushroom based solely on an image of it?

# The Approach


Collect a large set of images for each genus of mushroom we want to classify

For this, we leverage the existing Kaggle dataset:

***Mushrooms classification - Common genus's images***

Consists of common mushrooms from Northern Europe (should be applicable to US as well)





# Not all images are useful

- We cull certain images from dataset, such as:

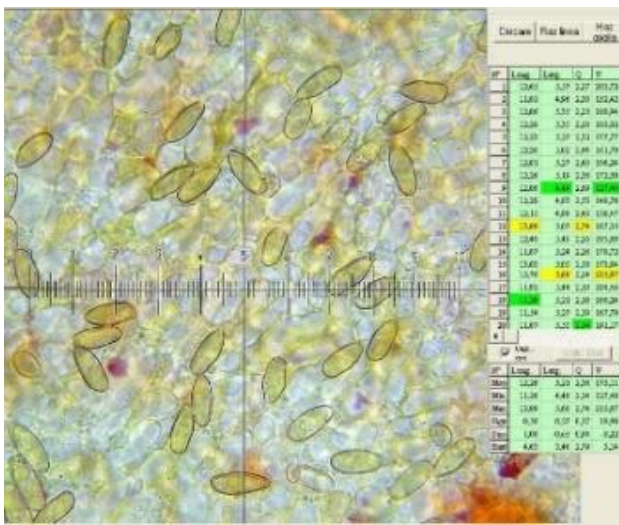


Arrows  
and text


Closeups of stems



Not mushrooms



Microscopic  
images



Extreme closeups




# Some we leave in

Non-natural  
backgrounds



Sliced Mushrooms



Reasonable Closeups



# Problems of Variety Within Genus

All of these  
mushrooms are from  
the *Hygrocybe* Genus





# Problems of Similarity Across Genera

With a variety of  
appearance,  
there can be  
some overlap

Boletus



Suillus



Lactarius



Russula



Amanita




Entoloma





# Making all Images Identical Inputs

- Wide range of sizes and aspect ratios
- Transform each image:
  - Center-crop
  - Resize each image to 300X300
  - Retain RGB channels



Original Image of a Hygrocybe



Cropped, then resized to 300 X 300



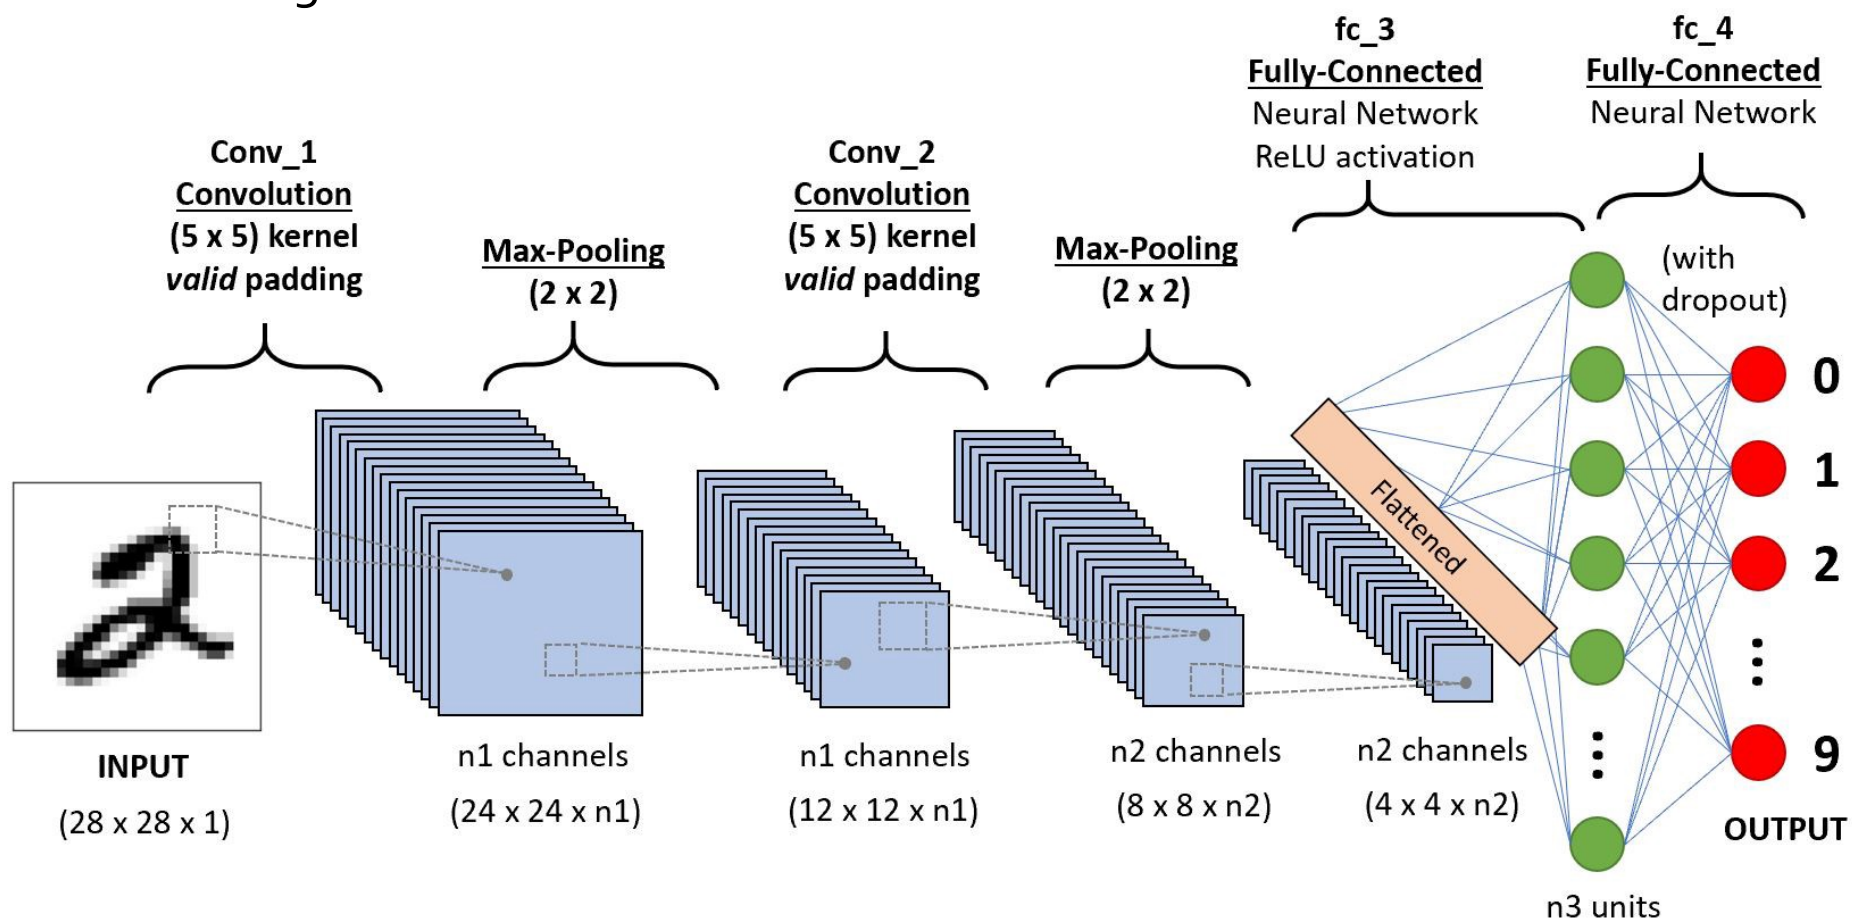


		Predicted Class	
		Default	PIF
True Class	Default	619	128
	PIF	95	3487

Modeling

# Convolutional Neural Network

- Breaks an image down into a series of patterns (“channels”)
- Uses Dropouts to prevent overfitting





# Convolutional Neural Network

- 131,000 Trainable parameters
- Performs very well with binary classification (e.g., a dataset consisting of only Boletus and Russula)

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 298, 298, 5)	140
=====		
max_pooling2d_1 (MaxPooling2D)	(None, 149, 149, 5)	0
=====		
conv2d_1 (Conv2D)	(None, 147, 147, 15)	690
=====		
max_pooling2d_2 (MaxPooling2D)	(None, 73, 73, 15)	0
=====		
conv2d_2 (Conv2D)	(None, 71, 71, 15)	2040
=====		
max_pooling2d_3 (MaxPooling2D)	(None, 35, 35, 15)	0
=====		
flatten (Flatten)	(None, 18375)	0
=====		
dropout (Dropout)	(None, 18375)	0
=====		
dense (Dense)	(None, 7)	128632
=====		

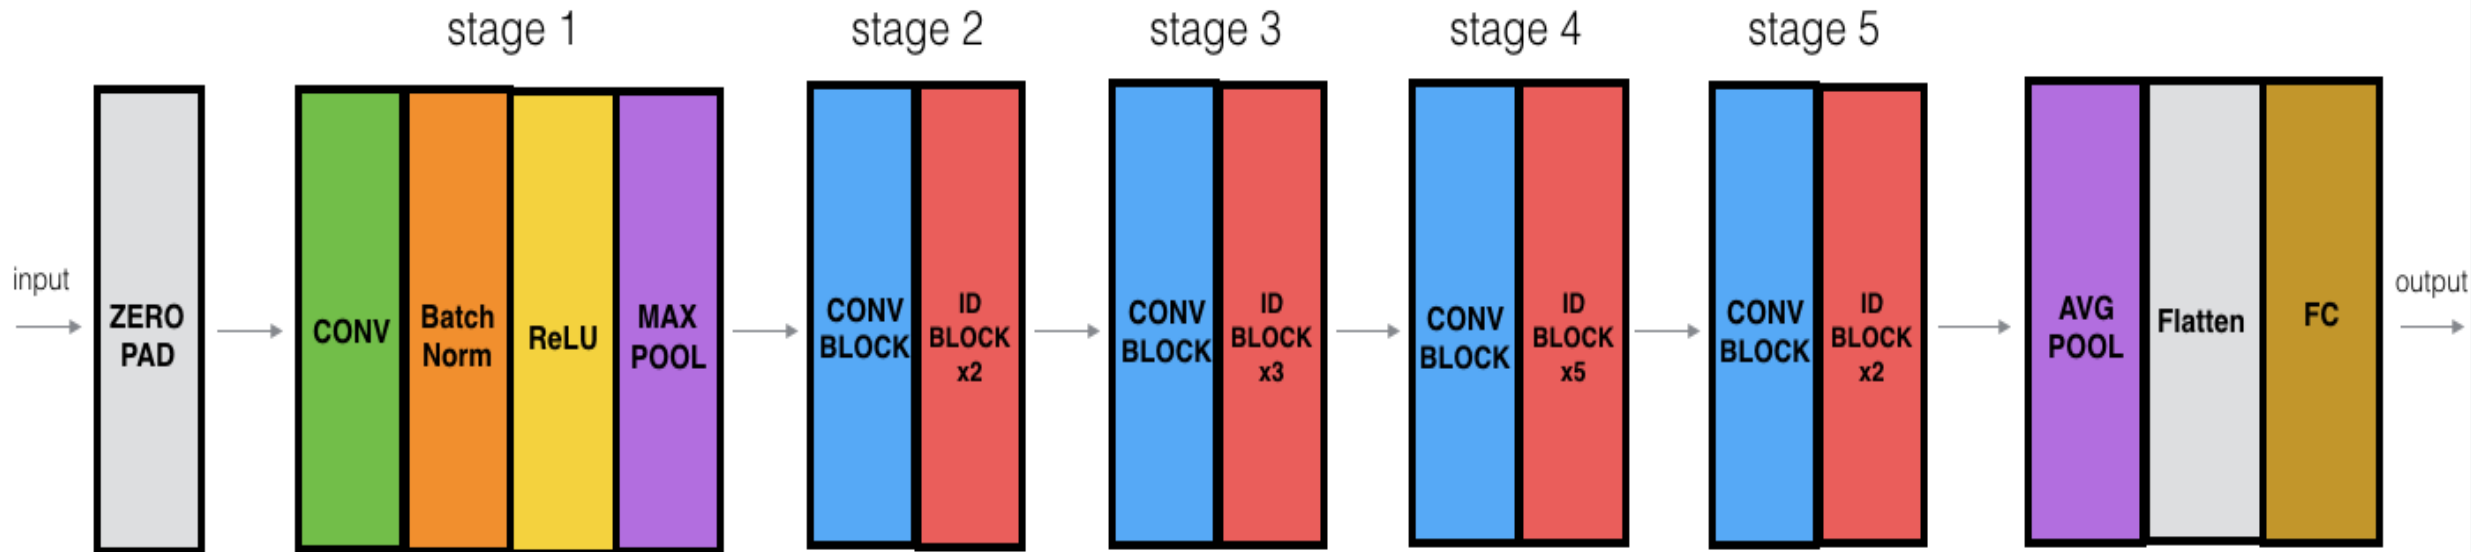
Total params: 131,502

Trainable params: 131,502

Non-trainable params: 0

# Residual Neural Network (ResNet50)

- Extremely deep learning (i.e., many layers) suffer from issues with vanishing gradients
- Skip connections across stages allow deep learning and greater classification accuracy
- ResNet50 comes pre-trained on ImageNet database (> 14 Million images)





# Residual Neural Network (ResNet50)

- Hold the ResNet50 layer as untrainable
- Still over half a million trainable parameters from Dense layers

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
=====	=====	=====
resnet50 (Model)	(None, 2048)	23587712
batch_normalization (BatchNo	(None, 2048)	8192
dense_1 (Dense)	(None, 256)	524544
batch_normalization_1 (Batch	(None, 256)	1024
dense_2 (Dense)	(None, 128)	32896
batch_normalization_2 (Batch	(None, 128)	512
dense_3 (Dense)	(None, 7)	903
=====	=====	=====

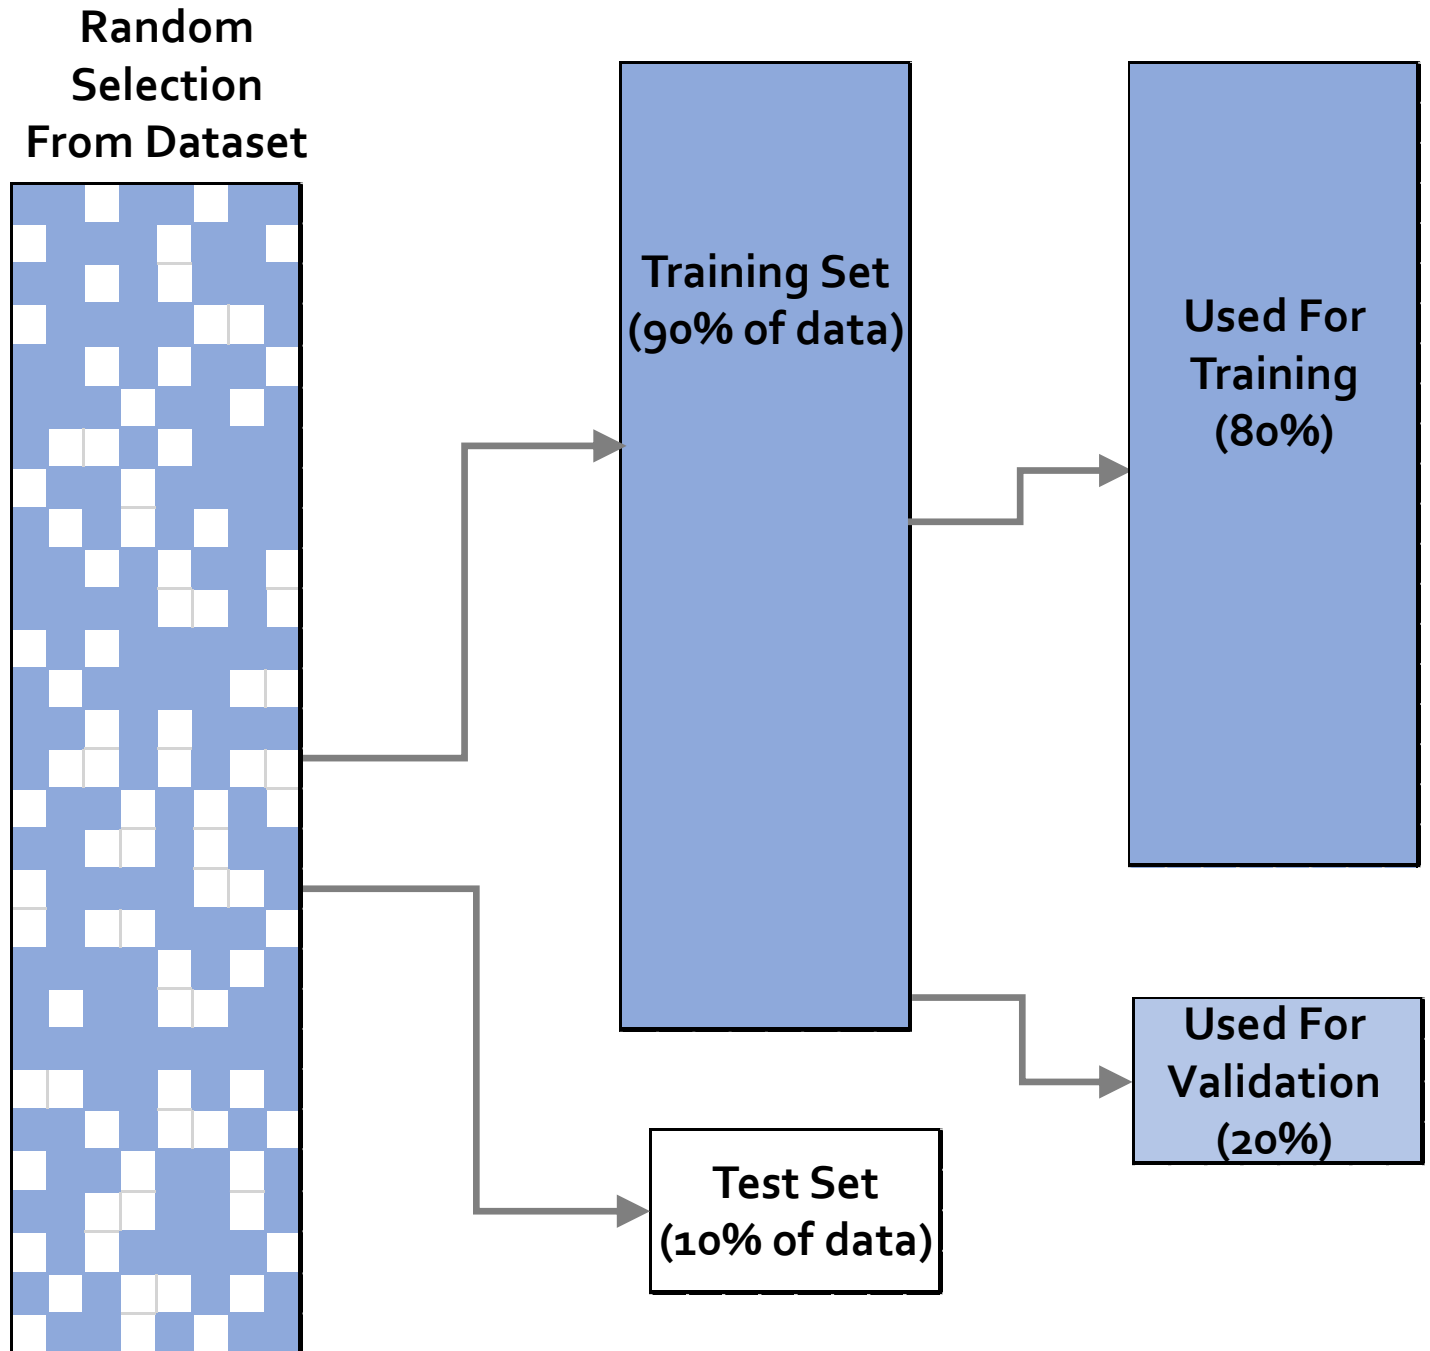
Total params: 24,155,783

Trainable params: 563,207


Non-trainable params: 23,592,576

# Classifier Models

- Nearly 5000 images
- Each image is a row in dataset, 300 X 300 X 3 long.
- Sci-kit Learn's Train/test split
- Training data further divided into training and validation sets
- Optimize model for Validation Set accuracy
- Test on Hold-Out Set




# Models Performance Comparison





# ResNet50 Model Performance: Confusion Matrix



# Model Performance Metrics

	TP	FP	TN	FN	Accuracy	Precision	Recall	F1-score
Hygrocybe	51.0	14.0	345.0	0.0	0.965854	0.784615	1.000000	0.879310
Boletus	69.0	1.0	338.0	2.0	0.992683	0.985714	0.971831	0.978723
Cortinarius	47.0	11.0	337.0	15.0	0.936585	0.810345	0.758065	0.783333
Russula	55.0	13.0	337.0	5.0	0.956098	0.808824	0.916667	0.859375
Suillus	44.0	3.0	350.0	13.0	0.960976	0.936170	0.771930	0.846154
Entoloma	44.0	4.0	349.0	13.0	0.958537	0.916667	0.771930	0.838095
Amanita	50.0	4.0	354.0	2.0	0.985366	0.925926	0.961538	0.943396

- False Positive: Thinking a mushroom is a certain type, when it's not
- False Negative: Thinking a mushroom is a different type, when it's not
- Best classification performance with Amanita and Boletes
- Worst classification performance overall with Cortinarius

# Examples of Misclassified Mushrooms

- The model gets confused, because mushrooms can be, well, confusing!
- Dirty, sliced, upside down: many of these are non-ideal presentations

Called Russula, Really Cortinarius



Called Amanita, Really Entoloma



Called Hygrocybe, Really Suillus



Called Russula, Really Cortinarius



Called Cortinarius, Really Suillus



Called Suillus, Really Entoloma



Called Entoloma, Really Russula



Called Cortinarius, Really Entoloma



Called Entoloma, Really Russula





# Conclusions

- ResNet50 model is way better at classifying mushrooms than a lay person such as myself. Aggregated accuracy across all seven mushroom genus is approaching 90%.
- Arguably on par with what a human mushroom expert could be expected to do based on images alone.
- Useful as a tool to identify edible mushrooms, such as Boletus (aka Porcini), with 99% accuracy.



# Thank You!

more info at [github repository](#)

